

# ALGORITHMIC SKELETONS FOR BRANCH & BOUND

Michael Poldner

*University of Münster, Department of Information Systems  
Leonardo Campus 3, D-48149 Münster, Germany  
Email: poldner@wi.uni-muenster.de*

Herbert Kuchen

*University of Münster, Department of Information Systems  
Leonardo Campus 3, D-48149 Münster, Germany  
Email: kuchen@uni-muenster.de*

**Keywords:** Parallel Computing, Algorithmic Skeletons, Branch & Bound, Load Distribution, Termination Detection.

**Abstract:** Algorithmic skeletons are predefined components for parallel programming. We will present a skeleton for branch & bound problems for MIMD machines with distributed memory. This skeleton is based on a distributed work pool. We discuss two variants, one with supply-driven work distribution and one with demand-driven work distribution. This approach is compared to a simple branch & bound skeleton with a centralized work pool, which has been used in a previous version of our skeleton library Muesli. Based on experimental results for two example applications, namely the  $n$ -puzzle and the traveling salesman problem, we show that the distributed work pool is clearly better and enables good runtimes and in particular scalability. Moreover, we discuss some implementation aspects such as termination detection as well as overlapping computation and communication.

## 1 INTRODUCTION

Today, parallel programming of MIMD machines with distributed memory is mostly based on message-passing libraries such as MPI (W. Gropp, 1999; MPI, 2006). The resulting low programming level is error-prone and time consuming. Thus, many approaches have been suggested, which provide a higher level of abstraction and an easier program development. One such approach is based on so-called *algorithmic skeletons* (Cole, 1989; Cole, 2006), i.e. typical patterns for parallel programming which are often offered to the user as higher-order functions. By providing application-specific parameters to these functions, the user can adapt an application independent skeleton to the considered parallel application. (S)he does not have to worry about low-level implementation details such as sending and receiving messages. Since the skeletons are efficiently implemented, the resulting parallel application can be almost as efficient as one based on low-level message passing.

Algorithmic skeletons can be roughly divided into data parallel and task parallel ones. Data-parallel skeletons (see e.g. (R. Bisseling, 2005; G. H. Botorog, 1996; G. H. Botorog, 1998; H. Kuchen, 1994; Kuchen, 2002; Kuchen, 2004)) process a distributed data structure such as a dis-

tributed array or matrix as a whole, e.g. by applying a function to every element or by rotating or permuting its elements. Task-parallel skeletons (A. Benoit, 2005; Cole, 2004; Hofstedt, 1998; H. Kuchen, 2002; Kuchen, 2002; Kuchen, 2004; Pelagatti, 2003) construct a system of processes communicating via streams of data. Such a system is mostly generated by nesting typical building blocks such as farms and pipelines. In the present paper, we will focus on a particular task-parallel skeleton, namely a branch & bound skeleton.

Branch & bound (G.L. Nemhauser, 1999) is a well-known and frequently applied approach to solve certain optimization problems, among them integer and mixed-integer linear optimization problems (G.L. Nemhauser, 1999) and the well-known traveling salesman problem (J.D.C. Little, 1963). Many practically important but NP-hard planning problems can be formulated as (mixed) integer optimization problems, e.g. production planning, crew scheduling, and vehicle routing. Branch & bound is often the only practically successful approach to solve these problems exactly. In the sequel we will assume without loss of generality that an optimization problem consists of finding a solution value which minimizes an objective function while observing a system of constraints. The main idea of branch & bound is the fol-

lowing. A problem is recursively divided into subproblems and lower bounds for the optimal solution of each subproblem are computed. If a solution of a (sub)problem is found, it is also a solution of the overall problem. Then, all other subproblems can be discarded, whose corresponding lower bounds are greater than the value of the solution. Subproblems with smaller lower bounds still have to be considered recursively.

Only little related work on algorithmic skeletons for branch & bound can be found in the literature (E. Alba, 2002; F. Almeida, 2001; I. Dorta, 2003; Hofstedt, 1998). However, in the corresponding literature there is no discussion of different designs. The MaLLBa implementation is based on a master/worker scheme and it uses a central queue (rather than a heap) for storing problems. The master distributes problems to workers and receives their solutions and generated subproblems. On a shared memory machine this approach can work well. We will show in the sequel that a master/worker approach is less suited to handle branch & bound problems on distributed memory machines. In a previous version of the Muesli skeleton library, a branch & bound skeleton with a centralized work pool has been used, too (H. Kuchen, 2002). Hofstedt outlines a B&B skeleton with a distributed work pool. Here, work is only shared, if a local work pool is empty. Thus, worthwhile problems are not propagated quickly and their investigation is concentrated on a few workers only.

The rest of this paper is structured as follows. In Section 2, we recall, how branch & bound algorithms can be used to solve optimization problems. In Section 3, we introduce different designs of branch & bound skeletons in the framework of the skeleton library Muesli (Kuchen, 2002; Kuchen, 2004; Kuchen, 2006). After describing the simple centralized design considered in (H. Kuchen, 2002), we will focus on a design with a distributed work pool. Section 4 contains experimental results demonstrating the strengths and weaknesses of the different designs. In Section 5, we conclude and point out future work.

## 2 BRANCH & BOUND

Branch & bound algorithms are general methods used for solving difficult combinatorial optimization problems. In this section, we illustrate the main principles of branch & bound algorithms using the 8-puzzle, a simplified version of the well-known 15-puzzle (Quinn, 1994), as example. A branch & bound algorithm searches the complete solution space of a given problem for the best solution. Due to the exponentially increasing number of feasible solutions, their explicit enumeration is often impossible in prac-

tice. However, the knowledge about the currently best solution, which is called *incumbent*, and the use of *bounds* for the function to be optimized enables the algorithm to search parts of the solution space only implicitly. During the solution process, a pool of yet unexplored subsets of the solution space, called the *work pool*, describes the current status of the search. Initially there is only one subset, namely the complete solution space, and the best solution found so far is infinity. The unexplored subsets are represented as nodes in a dynamically generated search tree, which initially only contains the root, and each iteration of the branch & bound algorithm processes one such node. This tree is called the *state-space tree*. Each node in the state-space tree has associated data, called its *description*, which can be used to determine, whether it represents a *solution* and whether it has any successors. A branch & bound problem is solved by applying a small set of basic rules. While the signature of these rules is always the same, the concrete formulation of the rules is problem dependent. Starting from a given initial problem, subproblems with pairwise disjoint state spaces are generated using an appropriate *branching rule*. A generated subproblem can be estimated applying a *bounding rule*. Using a *selection rule*, the subproblem to be branched from next is chosen from the work pool. Last but not least subproblems with non-optimal or inadmissible solutions can be eliminated during the computation using an *elimination rule*. The sequence of the application of these rules may vary according to the strategy chosen for selecting the next node to process (J. Clausen, 1999). As an example of the branch and bound technique, consider the 8-puzzle (Quinn, 1994). Figure 1 illustrates the goal state of the 8-puzzle and the first three levels of the state-space tree.

The 8-puzzle consists of eight tiles, numbered 1 through 8, arranged on a  $3 \times 3$  board. Eight positions on the board contain exactly one tile and the remaining position is empty. The objective of the puzzle is to repeatedly fill the hole with a tile adjacent to it in horizontal or vertical direction, until the tiles are in row major order. The aim is to solve the puzzle in the least number of moves.

The branching rule describes, how to split a problem represented by a given initial board into subproblems represented by the boards resulting after all valid moves. A minimum number of tile moves needed to solve the puzzle can be estimated by adding the number of tile moves made so far to the Manhattan distance between the current position of each tile and its goal position. The computation of this lower bound is described by the bounding rule.

The state-space tree represents all possible boards that can be reached from the initial board. One way to solve this puzzle is to pursue a breadth first search or a depth first search of the state-space tree until the




Figure 1: Upper part of the state-space tree corresponding to an instance of the 8-puzzle and its goal board.

sorted board is discovered. However, we can often reach the goal faster by selecting the node with the best lower bound to branch from. This selection rule corresponds to a best-first search strategy. Other selection rules such as a variant of depth-first search are discussed in (J. Clausen, 1999; Y. Shinano, 1995; Y. Shinano, 1997).

Branch & bound algorithms can be parallelized at a low or at a high level. In case of a low-level parallelization, the sequential algorithm is taken as a starting point and just the computation of the lower bound, the selection of the subproblem to branch from next, and/or the application of the elimination rule are performed by several processes in a data parallel way. The overall behavior of such a parallel algorithm resembles of the sequential algorithm.

In case of a high-level parallelization, the effects and consequences of the parallelism are not restricted to a particular part of the algorithm, but influence the algorithm as a whole. Several iterations of the main loop are performed in a task-parallel way, such that the state-space tree is explored in a different (non-deterministic!) order than in the sequential algorithm.

### 3 BRANCH & BOUND SKELETONS

In this section, we will consider different implementation and design issues of branch & bound skeletons. For the most interesting distributed design, several work distribution strategies are discussed and compared with respect to scalability, overhead, and performance. Moreover, a corresponding termination detection algorithm is presented.

A B&B skeleton is based on one or more branch & bound algorithms and offers them to the user as predefined parallel components. Parallel branch & bound algorithms can be classified depending on the organization of the work pool. A central, distributed,

and hybrid organization can be distinguished. In the MaLLBa project, a central work pool is used (F. Almeida, 2001; I. Dorta, 2003). Hofstedt (Hofstedt, 1998) sketches a distributed scheme, where work is only delegated, if a local work pool is empty. Shinano et al. (Y. Shinano, 1995; Y. Shinano, 1997) and Xu et al. (Y. Xu, 2005) describe hybrid approaches. A more detailed classification can be found in (Trienekens, 1990), where also complete and partial knowledge bases, different strategies for the use of knowledge and the division of work as well as the chosen synchronicity of processes are distinguished.

Moreover, different selection rules can be fixed. Here, we use the classical best-first strategy. Let us mention that this can be used to simulate other strategies such as the depth-first approach suggested by Clausen and Perregaard (J. Clausen, 1999). The bounding function just has to depend on the depth in the state-space tree.

We will consider the skeletons in the context of the skeleton library Muesli (Kuchen, 2002; Kuchen, 2004; Kuchen, 2006). Muesli is based on MPI (W. Gropp, 1999; MPI, 2006) internally in order to inherit its platform independence.

#### 3.1 Design with a Centralized Work Pool Manager

The simplest approach is a kind of the master/worker design as depicted in Figure 2. The work pool is maintained by the master, which distributes problems to the workers and receives solutions and subproblems from them. The approach taken in a previous version of the skeleton library Muesli is based on this centralized design. When a worker receives a problem, it either solves it or decomposes it into subproblems and computes a lower bound for each of the subproblems. The work pool is organized as a heap, and the subproblem with the best lower bound at the time is stored in its root. Idle workers are served with new problems taken from the root. This selection

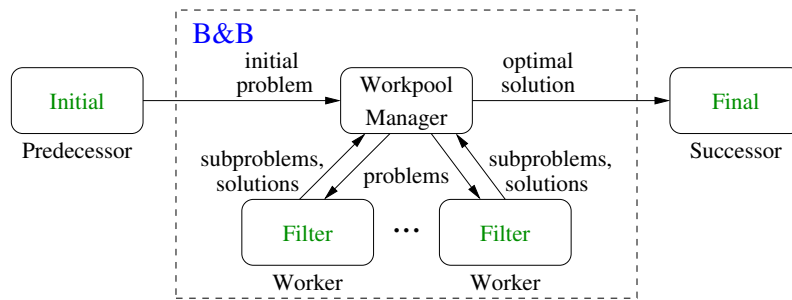


Figure 2: Branch &amp; bound skeleton with centralized work pool manager.

rule implicitly implements a best-first search strategy. Subproblems are discarded, if their bounds indicate that they cannot produce better solutions than the best known solution. An optimal solution is found, if the master has received a solution, which is better than all the bounds of all the problems in its work pool and no worker currently processes a subproblem. If at least one worker is processing, it can lead to a new incumbent. When the execution is finished, the optimal solution is sent to the master's successor in the overall process topology<sup>1</sup> and the skeleton is ready to accept and solve the next optimization problem. The code fragment in Fig. 3 illustrates the application of our skeleton in the context of the Muesli library. It constructs the process topology shown in Fig. 2.

```
int main(int argc, char* argv[]) {
    InitSkeletons(argc, argv);
    // step 1: create a process topology
    Initial<Problem> initial(generateProblem);
    Filter<Problem, Problem> filter(generateCases, 1);
    BranchAndBound<Problem> bnb(filter, n,
                                betterThan, isSolution);
    Final<Problem> final(fin);
    Pipe pipe(initial, bnb, final);
    // step 2: start process topology
    pipe.start();
    TerminateSkeletons();
}
```

Figure 3: Example application using a branch and bound skeleton with centralized work pool manager.

In a first step the process topology is created using C++ constructors. The process topology consists of an initial process, a branch & bound process, and a final process connected by a pipeline skeleton. The initial process is parameterized with a `generateProblem` method returning the initial optimization problem that is to be solved. The filter process represents a worker. The passed function `generateCases` describes, how to branch & bound subproblems. The

<sup>1</sup>Remember that task-parallel skeletons can be nested.

constructor `BranchAndBound` produces  $n$  copies of the worker and connects them to the internal work pool manager (which is not visible to the user). `bool betterThan(Problem x1, Problem x2)` has to deliver true, iff the lower (upper) bound for the best solution of problem  $x1$  is better than the lower (upper) bound for the best solution of problem  $x2$  in case of a minimization (maximization) problem. This function is used internally for the work pool organization. The function `bool isSolution(Problem x)` can be used to discover, whether its argument  $x$  is a solution or not. The final process receives and processes the optimal solution. Problems and solutions are encoded by the same type `Problem`.

The advantage of a single central work pool maintained by the master is that it provides a good overall picture of the work still to be done. This makes it easy to provide each worker with a good subproblem to branch from and to prune the work pool. Moreover, the termination of the workers is easy to implement, because the master knows about all idle workers at any time, and the best solution can be detected easily. The disadvantage is that accessing the work pool tends to be a bottleneck, as the work pool can only be accessed by one worker at a time. This may result in high idle times on the workers' site. Another disadvantage is that the master/worker approach incurs high communication costs, since each subproblem is sent from its producer to the master and propagated to its processing worker. If the master decides to eliminate a received subproblem, time is wasted for its transmission. Moreover, the communication time required to send a problem to a worker and to receive in return some subproblems may be greater than the time needed to do the computation locally. The master's limited memory capacity for maintaining the work pool is another disadvantage of this architecture.

As we will see in the next subsection, these disadvantages can be avoided by a distributed maintenance of the work pool. However, this design requires a suitable scheme for distributing subproblems and some distributed termination detection.

### 3.2 Distributed Work Pool

Figure 5 illustrates the design of the distributed branch and bound (DBB) skeleton provided by the Muesli skeleton library. It consists of a set of peer solvers, which exchange problems, solutions, and (possibly) load information. Several topologies for connecting the solvers are possible. For small numbers of processors, a ring topology can be used, since it enables an easy termination detection. For larger numbers of processors, topologies like torus or hypercube may lead to a faster propagation of work from hot spots to idle processors. For simplicity, we will assume a ring topology in the sequel. Compared to more complicated topologies the ring also simplifies the dynamic adaption of the number of workers in case that more or less computation capacity has to be devoted to the branch & bound skeleton within the overall computation. This (not yet implemented) feature will enable a well-balanced overall computation.

In our example,  $n = 5$  solvers are used. Each solver maintains its own local work pool and has one entrance and one exit. Exactly one of the solvers, called the *master solver*, serves as an entrance to the DBB-skeleton and receives new optimization problems from the predecessor. Any of the  $n$  solvers may deliver the detected optimal solution to the successor of the branch & bound skeleton in the overall process topology. All solvers know each other for a fast distribution of newly detected best solutions<sup>2</sup>. If the skeleton only consists of a single solver neither communication nor distributed termination detection are necessary. In this case all communication parts as well as the distributed termination detection algorithm are bypassed to speed up the computation.

The code fragment in Fig. 4 shows an example application of our distributed B&B skeleton. It constructs the process topology depicted in Fig. 5. Work request messages are only sent when using a demand-driven work distribution.

The construction of the process topology resembles that in the previous example. Instead of a filter a `BBSolver` process is used as a worker. In addition to the `betterThan` and `isSolution` function two other argument functions are passed to the constructor, namely a `branch` and a `bound` function. The constructor `DistributedBB` produces  $n$  copies of the solver. One of the solvers is automatically chosen as the master solver.

As described in the previous section, a task-parallel skeleton consumes a stream of input values and produces a stream of output values. If the master solver receives a new optimization problem, the communication with the predecessor is blocked until the received

<sup>2</sup>Thus, the topology is in fact a kind of wheel with spokes rather than a ring.

```
int main(int argc, char* argv) {
    InitSkeletons(argc,argv);
    // step 1: create a process topology
    Initial<Problem> initial(generateProblem);
    BBSolver<Problem> solver("ring",branch,bound,
                           betterThan,isSolution);
    DistributedBB<Problem> bnb =
        DistributedBB<Problem>(solver,n);
    Final<Problem> final(fin);
    Pipe pipe(initial,bnb,final);
    // step 2: start process topology
    pipe.start();
    TerminateSkeletons();
}
```

Figure 4: Task parallel example application of a fully distributed Branch and Bound skeleton.

problem is solved. This ensures that the skeleton processes only one optimization problem at a time. There are different variants for the initialization of parallel branch & bound algorithms with the objective of providing each worker with a certain amount of work within the start-up phase. Ideally, the work load is distributed equally to all workers. However, the work load is hard to predict without any domain knowledge. For this reason the skeleton uses the most common approach, namely *root initialization*, i.e. the root of the state space tree is inserted into the local work pool of the master solver. Subproblems are distributed according to the load balancing scheme applied by the solvers. This initialization has the advantage that it is very easy to implement and no additional code is necessary. Other initialization strategies are discussed in the literature. A good survey can be found in (Henrich, 1994a).

Each worker repeatedly executes two phases: a communication phase and a solution phase. Let us first consider the communication phase. In order to avoid that computation time is wasted with the solution of irrelevant subproblems, it is essential to spread and process new best solutions as quickly as possible. For this reason, we distinguish problem messages and incumbent messages. Each solver first checks for arriving incumbents with `MPL::Testsome`. If it has received new incumbents, the solver stores the best and discards the others. Moreover, it removes subproblems whose lower bound is worse than the incumbent from the work pool. Then, it checks for arriving subproblems and stores them in the work pool, if their lower bounds are better than the incumbent.

The solution phase starts with selecting an unexamined subproblem from the work pool. As in the master/worker design, the work pool is organized as a heap and the selection rule implements a best-first search strategy. The selected problem is decomposed into  $m$  subproblems by applying `branch`. For each



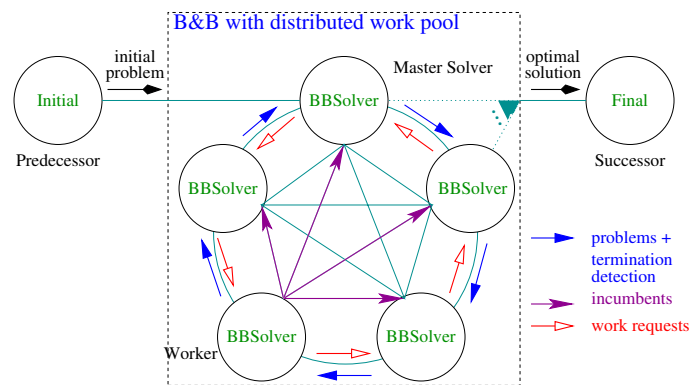


Figure 5: Branch &amp; bound skeleton with distributed work pool.

of the subproblems, we proceed as follows. First, we check, whether it is solved. If a new best solution is detected, we update the local incumbent and broadcast it. A worse solution is discarded. Finally, if the subproblem is not yet solved, the `bound` function is applied and the subproblem is stored in the work pool (see Fig. 5).

### 3.3 Load Distribution and Knowledge Sharing

Since the work pools of the different solvers, grow and shrink differently, some load balancing mechanism is required. Many global and local load distribution schemes have been studied in the literature (Henrich, 1994b; Henrich, 1995; R. Lüling, 1992; N. Mahapatra, 1998; Sanders, 1998; A. Shina, 1992) and many of them are suited in the context of a distributed branch & bound skeleton. Here, we will focus on two local load balancing schemes, a supply- and a demand-driven one. The local schemes avoid the larger overhead of a global scheme. On the other hand, they need more time to distribute work over long distances.

With the simple supply-driven scheme, each worker sends in each  $i$ th iteration its second best problem to its right neighbor in the ring topology. It always processes the best problem itself, in order to avoid communication overhead compared to the sequential algorithm. The supply driven approach has the advantage that it distributes work slightly more quickly than a demand driven approach, since there is no need for work requests. This may be beneficial in the beginning of the computation. A major disadvantage of this approach is that many subproblems are transmitted in vain, since they will be sooner or later discarded at their destination due to better incumbents, in particular for small  $i$ . Thus, high communication costs are caused.

The demand-driven approach distributes load only in case that a neighbor requests it. In our case, a neighbor sends the lower bound of the best problem in its work pool (see Fig. 5). If this value is worse than the lower bound of the second best problem of the worker receiving this information, it is interpreted as a work request and a problem is transmitted to the neighbor. In case that the work pool of the neighbor is empty, the information message indicates this fact rather than transmitting a lower bound. An information message is sent every  $i$ th iteration of the main loop. In order to avoid flooding the network with "empty work pool" messages, such messages are never sent twice. If the receiver of an "empty work pool message" is idle, too, it stores this request and serves it as soon as possible. The advantage of this algorithm is that distributing load only occurs, if it is necessary and beneficial. The overhead of sending load information messages is very low due to their small sizes. For small  $i$  the overhead is bigger, but idle processors get work more quickly.

### 3.4 Termination Detection

In the distributed setting, it is harder to detect that the computation has finished and the optimal solution has been found. The termination detection algorithm used in the DBB-skeleton is a variant of Dijkstra's algorithm outlined in (Quinn, 1994). Our implementation utilizes the specific property of MPI that the order in which messages are received from a sender  $S$  is always equal to the order in which they were sent by  $S$ . This characteristic can be used for the purpose of termination detection in connection with local load distribution strategies as described above.

As mentioned, we arrange the workers in a ring topology, since this renders the termination detection particularly easy and simplifies the dynamic addition and removal of workers. For a small number of processors (as in our system), the large diameter of the




Figure 6: Runtimes for the 16 city TSP using the central work pool manager and the distributed work pool with supply- and demand-driven work distribution depending on the number of workers.




Figure 7: Speedups for the 30 city TSP using the distributed work pool with demand-driven work distribution depending on the number of workers. The speedups are the averages taken from 300 runs with different, randomly generated maps.

ring topology is no serious problem for the distribution of work.

Let  $n$  be the number of solvers of the DBB-skeleton. When the master solver receives a new optimization problem, it initializes the termination detection by sending a token along the ring in the same direction as the load is distributed. The token only consists of an `int` value. Initially, the token has the value  $n$ . If a solver receives a new subproblem, this event is noted by setting a flag to `true`. On arrival of a token the solver uses the rules stated by the following pseudo code:

```

IF (workpool is empty AND flag == false)
    token := token - 1;
IF (workpool is not empty OR flag == true) {
    token := n; flag := false; }
IF (token > 0) send token to successor;
IF (token == 0) computation is finished;

```

Only if all workers are idle, the token is decremented by every worker and the computation is finished. No more problems can be in the network, since the token cannot overtake other messages on its way. Note that this algorithm only works for load balancing strate-

gies which send load in the same direction as the token.

## 4 EXPERIMENTAL RESULTS

We have tested the different versions of the branch & bound skeleton experimentally on a IBM workstation cluster (ZIV, 2006) using up to 16 Intel Xeon EM64T processors with 3.6 GHz, 1 MB L2 Cache, and 4 GB memory, connected by a Myrinet (Myricom, 2006). As example applications we have considered the  $n$ -puzzle as explained in section 2 as well as a parallel version of the traveling salesman problem (TSP) algorithm by Little et al. (J.D.C. Little, 1963). Both differ w.r.t. the quality of their bounding functions and hence in the number of considered irrelevant subproblems.

The presented B&B algorithm for the  $n$ -puzzle has a rather bad bounding function based on the Manhattan distance of each tile to its destination. It is bad, since the computed lower bounds are often much below the value of the best solution. As a consequence, the best-first search strategy is not very effective and the number of problems considered by the parallel skeleton differs enormously over several runs with the same inputs. This number largely depends on the fact whether a subproblem leading to the optimal solution is picked up early or late. Note that the parallel algorithm behaves non-deterministically in the way the search-space tree is explored. In order to get reliable results, we have repeated each run 100 times and computed the average runtimes.

The goal of the TSP is to find the shortest round trip through  $n$  cities. Little's algorithm represents each problem by its residual adjacency matrix, a set of chosen edges representing a partially completed tour, and a lower bound on the length of any full tour, which can be generated by extending the given partial tour. New problems are produced by selecting a *key edge* and generating two new problems, in which the chosen edge is included and excluded from the emerging tour, respectively. The key edge is selected based on the impact that the exclusion of the edge will have on the lower bound. The lower bounds are computed based on the fact that each city has to be entered and left once and that consequently one value in every row and column of the adjacency matrix has to be picked. The processing of a problem mainly requires three passes through the adjacency matrix.

The TSP algorithm computes rather precise lower bounds. Thus, the best-first strategy works fine, and the parallel implementation based on Quinn's formulation of the algorithm (Quinn, 1994) considers only very few problems more than the sequential algorithm, as explained below.




Figure 8: Speedups for 24-puzzle using the central work pool manager and the distributed work pool with demand-driven work distribution depending on the number of workers.

Table 1: Distribution of problems for the 16 city TSP using a distributed work pool and demand driven work distribution.

#workers	runtime (s)	# considered problems							
		total	worker						
1	10.38	263019	263019						
2	5.54	274002	139922	134080					
3	3.52	263583	90783	86039	86761				
4	2.64	262794	66536	65141	65475	65642			
5	2.10	273175	55863	52386	55993	53878	55055		
6	1.74	270525	45916	45150	44938	45574	42638	46309	
7	1.52	263180	39495	38749	37492	37197	37134	36273	36840
8	1.35	265698	34196	33763	33525	32793	32424	32466	32231

Consequently, the runtimes were relatively similar over several runs with the same parameters. For the TSP, we have used a real world 16 city map taken and adapted from (Reinelt, 1991) and 300 randomly generated 30 city maps. The real world map has much more sub-tours with similar lengths. Thus, proportionally more subproblems are processed which do not lead to the optimal solution than for the artificial map, where the best solution is found more easily.

When comparing the supply- and the demand-driven approach (see Figure 6 and the 3rd columns of Tables 1, 2), we notice that, as expected, the demand driven scheme is better, since it produces less communication overhead. The fact that the problems are distributed slightly slower causes no serious per-

Table 2: Distribution of problems for the 16 city TSP using a distributed work pool and supply driven work distribution.

#workers	runtime (s)	# considered problems							
		total	worker						
1	10.38	263019	263019						
2	5.84	262522	162536	99986					
3	3.92	271179	93060	89886	88233				
4	2.91	269021	66004	67709	66572	68736			
5	2.32	271717	53161	54569	55420	55074	53493		
6	2.03	265100	43227	47420	47739	42342	42185	42187	
7	1.75	265862	34390	34701	35100	36693	37917	51595	35466
8	1.35	264509	44379	32157	29789	29228	30017	31704	33124

Table 3: Distribution of problems for the 16 city TSP using a central work pool manager.

#workers	runtime (s)	# considered problems							
		total	worker						
1	22.01	263018	263018						
2	12.74	267057	133808	133249					
3	11.41	267019	116339	104349	46331				
4	11.26	267030	115396	103522	45945	2167			
5	11.26	267039	116064	103735	45406	1712	106		
6	11.25	267199	116082	103679	45470	1791	123	54	
7	11.26	267050	114111	103167	46558	2767	319	89	39
8	11.25	267024	115671	103675	45227	2071	226	81	45

formance penalty.

For the supply driven scheme, we have used an optimal number  $i$  for the amount of iterations that a worker waits before delegating a problem to a neighbor. If  $i$  is chosen too large, important problems will not spread out fast enough. If  $i$  is too small, the communication overhead will be too large. We found that the optimal value for  $i$  depends on the application problem and on the number of workers. If the number of workers increases,  $i$  has to be increased as well. In our experiments, the optimal values for  $i$  were ranging between 2 and 20 for up to 8 workers.

As expected, we see that for the centralized B&B skeleton the work pool manager quickly becomes a bottleneck and it has difficulties to keep more than 2 workers busy (see Figures 6, 8 and Table 3). This is due to the fact that the amount of computations done for a problem is linear in the size of the problem, just as the communication complexity for sending and receiving a problem. Thus, relatively little is gained by delegating a problem to a worker. The work pool manager has to spend only little work less for transmitting the problem than its processing would require. This property is typical for virtually all practically relevant branch & bound problems we are aware of. It has the important consequence that a centralized work pool manager does not work well for branch & bound on distributed memory machines. Also note that the centralized scheme needs one more processor, the work pool manager, than the distributed one rendering this approach even less attractive.

Both variants of the design with a distributed work pool do not have these drawbacks (see Figures 6, 7, 8 and Tables 1, 2). Here, the communication overhead is much smaller. Each worker fetches most problems from its own work pool, such that they require no communication. This is particularly true for the demand driven approach. This scheme has the advantage that after some start-up phase, in which all workers are supplied with problems, there is relatively little communication and the workers mainly process locally available problems. This is essential for achieving good runtimes and speedups. We anticipate that this insight not only applies to branch & bound but also to other skeletons with a similar characteristic



such as divide & conquer and other search skeletons. We are currently working on experimental results supporting this claim.

Interestingly we could even observe slightly super-linear speedups for the 30 city TSPs. They can be explained by the fact that a parallel B&B algorithm may tackle important subproblems earlier than the sequential one, since it processes the state-space tree in a different order (T. Lai, 1984).

It is clear that a parallel B&B algorithm will typically consider more problems than a corresponding sequential one, since it eagerly processes several problems in parallel, which would be discarded in the sequential case, since their lower bounds are worse than a detected solution. Interestingly for both considered example applications, TSP and  $n$ -puzzle, the corresponding overhead was very small and only few additional problems have been processed by the parallel implementation (see the 3rd columns of Tables 1, 2, 3). For instance, for the 16 city TSP no more than  $274002 - 263019 = 10983$  additional problems are processed by the parallel algorithm; this is less than 4.2 %. This is essential for achieving reasonable speedups.

As an implementation detail of the centralized approach let us mention that it is important that the work pool manager receives in each iteration all available subproblems and solutions from the workers rather than just one of them. The reason is that MPI\_Waitany (used internally) is unfair and that an overloaded work pool manager will hence almost exclusively communicate with a small number of workers. If a starving worker has an important subproblem (one that leads to the optimal solution) or a good solution, which it is not able to deliver to the work pool manager, this will cause very bad runtimes.

Another implementation detail of the centralized approach concerns the amount of buffering. In order to be able to overlap computation and communication, it is a good idea that the work pool manager not only sends one problem to each worker and then waits for the results, but that it sends  $m$  problems such that the worker can directly tackle the next problem after finishing the previous one. Here it turned out that one has to be careful not to choose  $m$  too large, since then problems which would otherwise be discarded due to appearing better incumbents will be processed (in vain). In our experiments,  $m = 2$  was a good choice.

## 5 CONCLUSION

We have considered two different implementation schemes for the branch & bound skeleton. Besides a simple approach with a central work pool manager,

we have investigated a scheme with a distributed work pool. As our analysis and experimental results show, the communication overhead is high for the centralized approach and the work pool manager quickly becomes a bottleneck, in particular, if the number of computation steps for each problem grows linearly with the problem size, as it is the case for virtually all practically relevant branch & bound problems. Thus, the centralized scheme does not work well in practice.

On the other hand, our scheme with a distributed work pool works fine and provides good runtimes and scalability. The latter is not trivial, as discussed e.g. in the book of Quinn (Quinn, 1994), since parallel B&B algorithms tend to process an increasing number of irrelevant problems the more processors are employed. In particular, the demand-driven design works well due to its low communication overhead.

For the supply-driven approach, we have investigated, how often a problem should be propagated to a neighbor. Depending on the application and the number of workers, we have observed the best runtimes, if a problem was delegated between every 2nd and every 20th iteration.

We are not aware of any previous comparison of different implementation schemes of branch & bound skeletons for MIMD machines with distributed memory in the literature. In the MaLLBa project (E. Alba, 2002; F. Almeida, 2001), a branch & bound skeleton based on a master/worker approach and a queue for storing subproblems has been developed. But as we pointed out above, this scheme is more suitable for shared memory machines than for distributed memory machines. Hofstedt (Hofstedt, 1998) sketches a B&B skeleton with a distributed work pool. Here, work is only delegated, if a local work pool is empty. A quick propagation of “interesting” subproblems are missing. According to our experience, this leads to a suboptimal behavior. Moreover, Hofstedt gives only few experimental results based on reduction steps in a functional programming setting rather than actual runtimes and speedups.

As future work, we intend to investigate alternative implementation schemes of skeletons for other search algorithms and for divide & conquer.

## REFERENCES

- A. Benoit, M. Cole, J. H. S. G. (2005). Flexible skeletal programming with eskel. In *Proc. EuroPar 2005*. LNCS 3648, 761–770, Springer Verlag, 2005.
- A. Shina, L. K. (1992). A load balancing strategy for prioritized execution of tasks. In *Proc. Workshop on Dynamic Object Placement and Load Balancing, ECOOP’92*.
- Cole, M. (1989). *Algorithmic Skeletons: Structured Management of Parallel Computation*. MIT Press.

- Cole, M. (2004). Bringing skeletons out of the closet: A pragmatic manifesto for skeletal parallel programming. In *Parallel Computing* 30(3), 389–406.
- Cole, M. (2006). The skeletal parallelism web page. <http://homepages.inf.ed.ac.uk/mic/Skeletons/>.
- E. Alba, F. Almeida, e. a. (2002). Mallba: A library of skeletons for combinatorial search. In *Proc. Euro-Par 2002*. LNCS 2400, 927–932, Springer Verlag, 2005.
- F. Almeida, I. Dorta, e. a. (2001). Mallba: Branch and bound paradigm. In *Technical Report DT-01-2*. University of La Laguna, Spain, Dpto. Estadística, I.O. y Computación.
- G. H. Botorog, H. K. (1996). Efficient parallel programming with algorithmic skeletons. In *Proc. Euro-Par'96*. LNCS 1123, 718–731, Springer Verlag, 1996.
- G. H. Botorog, H. K. (1998). Efficient high-level parallel programming. In *Theoretical Computer Science* 196, 71–107.
- G.L. Nemhauser, L. W. (1999). Integer and combinatorial optimization. Wiley.
- H. Kuchen, R. Plasmeijer, H. S. (1994). Efficient distributed memory implementation of a data parallel functional language. In *Proc. PARLE'94*. LNCS 817, 466–475, Springer Verlag.
- H. Kuchen, M. C. (2002). The integration of task and data parallel skeletons. In *Parallel Processing Letters* 12(2), 141–155.
- Henrich, D. (1994a). Initialization of parallel branch-and-bound algorithms. In *Proc. 2nd International Workshop on Parallel Processing for Artificial Intelligence (PPAI-93)*. Elsevier.
- Henrich, D. (1994b). Local load balancing for data-parallel branch-and-bound. In *Proc. Massively Parallel Processing Applications and Development*, 227–234.
- Henrich, D. (1995). Lastverteilung fuer feinkoernig parallelisiertes branch-and-bound. In *PhD Thesis*. TH Karlsruhe.
- Hofstedt, P. (1998). Task parallel skeletons for irregularly structured problems. In *Proc. EuroPar'98*. LNCS 1470, 676 – 681, Springer Verlag.
- I. Dorta, C. Leon, C. R. A. R. (2003). Parallel skeletons for divide and conquer and branch and bound techniques. In *Proc. 11th Euromicro Conference on Parallel, Distributed and Network-based Processing (PDP2003)*.
- J. Clausen, M. P. (1999). On the best search strategy in parallel branch-and-bound: Best-first search versus lazy depth-first search. In *Annals of Operations Research* 90, 1–17.
- J.D.C. Little, K.G. Murty, D. S. C. K. (1963). An algorithm for the traveling salesman problem. In *Operations Research* 11, 972–989.
- Kuchen, H. (2002). A skeleton library. In *Euro-Par'02*. LNCS 2400, 620–629, Springer Verlag.
- Kuchen, H. (2004). Optimizing sequences of skeleton calls. In *Domain-Specific Program Generation*. LNCS 3016, 254–273, Springer Verlag.
- Kuchen, H. (2006). The skeleton library web pages. <http://www.wi.uni-muenster.de/PI/forschung/Skeletons/index.php>.
- MPI (2006). Message passing interface forum, mpi. In *MPI: A Message-Passing Interface Standard*. <http://www.mpi-forum.org/docs/mpi-1.1-html/mpi-report.html>.
- Myricom (2006). The myricom homepage. <http://www.myri.com/>.
- N. Mahapatra, S. D. (1998). Adaptive quality equalizing: High-performance load balancing for parallel branch-and-bound across applications and computing systems. In *Proc. International Parallel Processing and Distributed Processing Symposium (IPDPS98)*.
- Pelagatti, S. (2003). Task and data parallelism in p3l. In *Patterns and Skeletons for Parallel and Distributed Computing*. eds. F.A. Rabhi, S. Gorlatch, 155–186, Springer Verlag.
- Quinn, M. (1994). *Parallel Computing: Theory and Practice*. McGraw Hill.
- R. Bisseling, I. F. (2005). Mondriaan sparse matrix partitioning for attacking cryptosystems – a case study. In *to appear in Proceedings of ParCo 2005, Malaga*.
- R. Lüling, B. M. (1992). Load balancing for distributed branch and bound algorithms. In *Proc. 6th International Parallel Processing Symposium (IPPS92)*, 543–549. IEEE.
- Reinelt, G. (1991). Tsplib – a traveling salesman problem library. In *ORSA Journal on Computing* 3, 376–384. see also: [http://www.iwr.uni-heidelberg.de/groups/comopt/software/TSPLIB95/\(gr17\)](http://www.iwr.uni-heidelberg.de/groups/comopt/software/TSPLIB95/(gr17)).
- Sanders, P. (1998). Tree shaped computations as a model for parallel applications. In *Proc. Workshop on Application Based Load Balancing (ALV'98)*. TU Munich.
- T. Lai, S. S. (1984). Anomalies in parallel branch-and-bound algorithms. In *Communications of the ACM* 27, 594–602.
- Trienekens, H. (1990). Parallel branch & bound algorithms. In *PhD Thesis*. University of Rotterdam.
- W. Gropp, E. Lusk, A. S. (1999). *Using MPI*. MIT Press.
- Y. Shinano, M. Higaki, R. H. (1995). A generalized utility for parallel branch and bound algorithms. In *Proc. 7th IEEE Symposium on Parallel and Distributed Processing*, 392–401. IEEE.
- Y. Shinano, M. Higaki, R. H. (1997). Control schemes in a generalized utility for parallel branch and bound algorithms. In *Proc. 11th International Parallel Processing Symposium*, 621–627. IEEE.
- Y. Xu, T. Ralphs, L. L. M. S. (2005). Alps: A framework for implementing parallel tree search algorithms. In *Proc. 9th INFORMS Computing Society Conference*.
- ZIV (2006). Ziv-cluster. <http://zivcluster.uni-muenster.de/>.